

Ex No: 6

BUILD A RECURRENT NEURAL NETWORK

Aim:

To build a recurrent neural network with Keras/TensorFlow.

Procedure:

1. Import required libraries for numerical computation and deep learning, such as NumPy and TensorFlow.
2. Define a function to generate dummy sequential data with random features and one-hot encoded labels.
3. Generate training data with a specified number of samples, sequence length, and output classes using the data generation function.
4. Build a sequential model using Keras with a SimpleRNN layer for processing sequential data and a Dense output layer for classification.
5. Compile the model with the Adam optimizer, categorical cross-entropy loss, and accuracy as a performance metric.
6. Train the model on the generated training data for a specified number of epochs and batch size.
7. Generate a separate set of sequential test data to evaluate the trained model's performance.
8. Evaluate the model using the test dataset and print the loss and accuracy values.
9. Generate new random sequential data for prediction.
10. Use the trained model to predict the class probabilities for the new data samples and print the predictions.

Code:

```
import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
# Generate dummy sequential data
def generate_data(num_samples=1000, sequence_length=10, num_classes=2):
    X = np.random.rand(num_samples, sequence_length, 1) # 3D input for RNN
```

```
y = np.random.randint(num_classes, size=(num_samples, num_classes)) # One-hot encoded labels

return X, y

X, y = generate_data()

model = keras.Sequential([
    layers.SimpleRNN(50, activation='relu', input_shape=(X.shape[1], X.shape[2])),
    layers.Dense(2, activation='softmax') # Adjust the number of units for your classes
])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Train the model

model.fit(X, y, epochs=10, batch_size=32)

# Generate test data

X_test, y_test = generate_data(num_samples=200)

# Evaluate the model

loss, accuracy = model.evaluate(X_test, y_test)

print(f'Test Accuracy: {accuracy:.4f}')

# Generate some new data for prediction

X_new = np.random.rand(5, 10, 1) # 5 new samples

predictions = model.predict(X_new)

print(predictions)
```

Output:

Epoch 1/10

```
C:\Users\mannu\anaconda3\envs\dlc\Lib\site-packages\keras\src\layers\rnn\rnn.py:204: UserWarning: Do
layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the
super().__init__(**kwargs)
```

32/32 ————— 2s 3ms/step - accuracy: 0.3543 - loss: 0.7056

Epoch 2/10

32/32 ————— 0s 3ms/step - accuracy: 0.3662 - loss: 0.7122

Epoch 3/10

32/32 ————— 0s 3ms/step - accuracy: 0.3485 - loss: 0.6976

Epoch 4/10

32/32 ————— 0s 4ms/step - accuracy: 0.2966 - loss: 0.7087

Epoch 5/10

32/32 ————— 0s 4ms/step - accuracy: 0.6222 - loss: 0.7392

Epoch 6/10

32/32 ————— 0s 3ms/step - accuracy: 0.6153 - loss: 0.7082

Epoch 7/10

32/32 ————— 0s 2ms/step - accuracy: 0.3386 - loss: 0.7176

Epoch 8/10

32/32 ————— 0s 2ms/step - accuracy: 0.3980 - loss: 0.7233

Epoch 9/10

32/32 ————— 0s 2ms/step - accuracy: 0.3866 - loss: 0.7133

Epoch 10/10

32/32 ————— 0s 2ms/step - accuracy: 0.4784 - loss: 0.6931

7/7 ————— 0s 3ms/step - accuracy: 0.4753 - loss: 0.6428

Test Accuracy: 0.4900

1/1 ————— 0s 200ms/step

[[0.47517636 0.52482367]

[0.504462 0.495538]

[0.4892657 0.51073426]

[0.49312642 0.50687355]

[0.5224694 0.47753054]]

Result:

Thus the program to build a recurrent neural network with Keras/TensorFlow has been executed successfully.